



# Debugging System Hangs on Solaris

Oliver Yang

Jun 25, 2007

Sun Microsystem

<http://oliveryang.net>

# Agenda

- What is a system hang?
- Debugging system hangs with kmdb/mdb
- Case Analysis
- References
- Appendix

# Agenda

- What is a system hang?

# What is a system hang?

- System...
  - > has no response
  - > is no longer usable
- What conditions cause hangs?
  - > Deadlock
  - > Resources exhaustion
  - > Hardware problems

# When system hangs happen...

- What you should do?
  - > Try a ping - reachable?
  - > Access network services - ssh/rsh/telnet... ?
  - > See console messages - any err/warn/fail ?
  - > System console status - no response?
  - > Try to force a system crash dump
  - > Check system logs and test journals after system boot

# Agenda

- What is a system hang?
- **Debugging system hangs with kmdb/mdb**
  - > Loading kmdb
  - > Forcing a crash dump
  - > Live debugging
  - > Crash dump analysis
  - > Other crash dump analysis tools

# Loading kmdb

- Boot-time Loading
  - > x86 - grub  
kernel /platform/i86pc/multiboot-k
  - > SPARC - OBP  
boot -k
- Runtime Loading
  - > mdb -K

# Forcing a crash dump

- General hangs
  - > Drop into kmdb or OBP
    - Keyboard - Stop+A or F1+A
    - Remote Console - Send a BREAK
  - > \$<systemdump (all platforms, if kmdb loaded)
  - > sync (OBP, SPARC only)
- Hard hangs
  - > You can't enter kmdb when running into hard hangs
  - > Enable deadman timer in /etc/system
    - set snooping=1



# Deadman panic

- Why panic?
  - > Setting in /etc/system file  
set snooping=1
  - > Deadman timer will trigger a panic if clock interrupt was inactive about 5000 ticks (1tick=10ms)
  - > We can change the default timer to other vaules(eg, 90s)  
set snoop\_interval = 90000000
- What we should do?
  - > Find out why clock interrupt become inactive

# Live debugging

- For special circumstance...
  - > Can't save crash dump
  - > System hangs occurred during system boot.
- System hang happened during boot
  - > Boot kmdb with the -kd options
  - > Set necessary variables for debugging:  
moddebug/W 0x80000000  
snooping/W 0x1  
kmem\_flags/W 0xf
  - > Set break point  
Using fully qualified symbol name - ::bp bge`bge\_attach

# Crash dump analysis

- Using mdb/kmdb and reading relevant source code
  - > To identify the set of kernel threads in deadlock
  - > To investigate how system hangs took place
- Checking crash dump files...
  - > System status checking
  - > Kernel threads checking
    - > CPU and dispatch queue
    - > User processes and kernel threads status
    - > Check the stack trace of suspicious threads:
      - Function name related to  
mutex(9F)/rwlock(9F)/condvar(9F)/semaphore(9F)/biowait(9F)
      - Running into an infinite loop

# **mdb - frequently used ::dcmds**

- System status checking
  - > System messages - ::msgbuf
  - > Clock interrupt - ::cycinfo
  - > Physical memory - ::memstat
  - > Cache/vmem allocation - ::kmastat
  - > Checking any necessary global variables
    - kmem\_flags/X
    - snooping/X
    - ...

# **mdb - frequently used ::dcmds**

- Kernel threads checking
  - > ::cpuinfo
  - > ::threadlist
  - > ::thread
  - > ::findstack
  - > ::mutex
  - > ::rwlock
  - > ::wchaninfo
  - > ::whatthread(Nevada only) or ::kgrep

# Other crash dump analysis tools

- ACT - Automated Crash Tool
  - > A complete list of threads with function arguments
  - > Detailed system setting and resource summary
  - > Deadlock detection - mutex and rwlocks only
  - > Threads blocked in either getblk() or biowait()
- SCAT - Solaris Crash Analysis Tool
- Download from <http://sunsolve.sun.com>

# Agenda

- What is a system hang?
- Debugging system hangs with kmdb/mdb
- **Case Analysis**
  - > Reverse locking order
  - > Infinite loop
  - > The constraints of current context

# Reverse locking order

- Multiple threads deadlocks on multiple locks
  - > To avoid the dead lock
    - > Must always lock in the same order
    - > Must always release in reverse order of locking
  - > Two threads acquiring two locks with reverse order
    - > Thread 1 ---> acquire Lock A ---> sleep and wait for Lock B
    - > Thread 2 ---> acquire Lock B ---> sleep and wait for Lock A
  - > See an example:
    - > <http://blog.ccw.com.cn/blog-htm-do-showone-itemid-12139-type-blog.html>



# Infinite loop - kmdb debugging

- Threads running into an infinite loop

```
>::cpuinfo -v
```

ID	ADDR	FLG	NRUN	BSPL	PRI	RNRN	KRNRN	SWITCH	THREAD	PROC
0	0000180c000	1d	0	0	0	yes	no	t-301817	30002e6e380	ifconfig

|  
 RUNNING <--  
 QUIESCED  
 EXISTS  
 ENABLE

```
> 30002e6e380::findstack -v
```

```
stack pointer for thread 30002e6e380: 2a101d04391
```

```

000002a101d04431 i_mod_hash_find_nosync+0x34(3000090bb40, 600042e5b58, 2a101d04e40, 1, 3000090bbb8, 98)
000002a101d044e1 mod_hash_find+0x18(3000090bb40, 600042e5b58, 2a101d04e40, 53, 30002e6e384, 0)
000002a101d04591 mac_open+0xf8(600042e5b58, 1, 600042e5c60, 70051800, 70057c00, 0)
000002a101d04751 dls_mac_hold+0x24(600042e5b58, 2a101d05350, 2a101d051c8, 1, 600042e5cc0, 0)
  
```

# Infinite loop - look into the code

- The infinite loop between L239 and L250

```
196 int
```

```
197 mac_open(const char *macname, uint_t ddi_instance, mac_handle_t *mhp)
```

```
198 {
```

```
...
```

```
239 again:
```

```
240     rw_enter(&i_mac_impl_lock, RW_WRITER);
```

```
...
```

```
248     if (mip->mi_destroying) {
```

```
249         rw_exit(&i_mac_impl_lock);
```

```
250         goto again;
```

```
251     }
```

# The constraints of current context

- Know the constraints of current context...
  - > Learn the constraints of current context. For example, high-level interrupt, interrupts threads, soft-interrupts, timeout(9F), etc...
  - > The lock usage should follow the constraints of current context, sometimes you have to use taskq(9F) when you can't block in current context
- The constraints of interrupt context
  - > Only spin mutex used in high-level interrupt(PIL>10)
  - > Interrupt thread(PIL<10) is a special context, be careful to use following functions in an interrupt context  
mutex(9F)/rwlock(9F)/condvar(9F)/semaphore(9F)/biowait(9F)

# The constraints of current context

- Self-Deadlock in an interrupt context
  - > An simple example:
    - > Drivers must call biodone(9F) when the transfer is complete to notify the thread blocked by biowait(9F). biodone(9F) is usually called in the interrupt routine.
    - > Call biowait(9F) in an interrupt thread
    - > But biodone(9F) only is called by the same interrupt thread

# The constraints of current context

- The implicit constraints of timeout(9F)
  - > See my blog: <http://blog.csdn.net/yayong>
  - > The clock interrupt thread trigger a taskq thread. That taskq thread ran into callout\_execute() by holding the mutex of callout table, then called the driver timeout handler registered by timeout(9F) routine
  - > The driver timeout handler blocked on a rwlock waiting for the rwlock owner releasing the rwlock
  - > At the same time, the owner of that rwlock was blocked on cv\_timedwait, and it has to wait for following clock interrupt to wake up it by calling callout\_execute()
  - > But the following callout\_execute() can't be done unless the previous callout\_execute() completed.

# Agenda

- What is a system hang?
- Debugging system hangs with kmdb/mdb
- Case Analysis
- **References**
  - > Books
  - > Blogs

# Books

- Solaris x86 Crash dump Analysis
  - > By Frank.Hofmann (2003-2005)
  - > Free download with Creative Commons Deed
  - > <http://opensolaris.org/os/community/documentation/files/book.pdf>
- Solaris Internals
  - > By Jim Mauro, Richard McDougall and Brendan Gregg
  - > 2nd Edition (July 10, 2006, ISBN 0131482092)
  - > Chinese edition have been published
  - > [http://www.solarisinternals.com/wiki/index.php/Solaris Internals](http://www.solarisinternals.com/wiki/index.php/Solaris_Internals)

# Blogs

- Kernel debugging part 1 kmdb
  - > [http://blogs.sun.com/eschrock/entry/kernel\\_debugging\\_part\\_1\\_kmdb](http://blogs.sun.com/eschrock/entry/kernel_debugging_part_1_kmdb)
- Debugging Solaris scheduling problems
  - > [http://blogs.sun.com/esaxe/entry/debugging\\_solaris\\_scheduling\\_problems\\_and](http://blogs.sun.com/esaxe/entry/debugging_solaris_scheduling_problems_and)
- A R/W deadlock of aggregation in GLD code
  - > <http://blog.ccw.com.cn/blog-htm-do-showone-itemid-12139-type-blog.html>
- Solaris learning journal(6)
  - > <http://blog.csdn.net/yayong/archive/2007/03/04/1520604.aspx>



# Agenda

- What is a system hang?
- Debugging system hangs with kmdb/mdb
- Case Analysis
- References
- **Appendix**
  - > Crash dump basics
  - > Modular Debugger – mdb(1)

# Crash dump basics

- It's similar with application core dumps, but...
  - > Dump system wide pages whereas a application core dump just contains the 1 specific process's pages
  - > It can be debugged by mdb whereas application code dumps can be debugged by dbx/gdb/mdb
  - > Managed by dumpadm(1M) whereas application code dumps are managed by coreadm(1M)
- dumpadm(1M)
  - > Dump content - kernel pages by default
  - > Dump device - /dev/dsk/c0t0d0s1 (swap)
  - > Savecore directory - /var/crash/<hostname>

# Crash dump basics

- `savecore(1M)`
  - > When the system is rebooted, `savecore` can be run to retrieve the image from the dump device and archive it to a disk file
- About crash dump files...
  - > `unix.X` - Symbol tables
  - > `vmunix.X` – Memory dump
  - > `bounds` - contains the sequence number to use for the next execution of `savecore`
  - > Check with `mdb X`
    - > eg. `mdb 0`

# Modular Debugger - mdb(1)

- mdb(1) basics
  - > commands (dcmd)
    - > ::dcmds for a list
    - > expression::dcmd - eg: cbd7bad8::ps
    - > ::help ::dcmd - ::help ::ps
  - > walkers
    - > ::walkers for a list
    - > expression::walk <walker\_name> - e.g. ::walk cpu
  - > macros
    - > \$M for a list
    - > \$<threadlist

# Modular Debugger - mdb(1)

- Symbols and typed data
  - > address::print (for symbol)
  - > address::print <type>
    - > eg. <address>::print cpu\_t
    - > eg. ::sizeof cpu\_t
- Pipelines
  - > expression, dcmd or walk can be piped
    - > ::walk <walk\_name> | ::dcmd
    - > e.g. ::walk cpu | ::print cpu\_t
  - > dcmd or walk can be piped with shell (mdb only)
    - > eg. ::ps ! grep bash



# Q &A

Oliver.Yang

Oliver.Yang@sun.com

<http://oliveryang.net>